

Supporting Information

MatterMind: An LLM-Assisted Platform for Crystal Generation, Materials Simulation, and AI-Guided Analysis

Yujun Lu^a, Kaiyang Xu^b and Lei Zhang^{a,*}

^aDepartment of Software Engineering, School of Software, Nanjing University of Information Science & Technology, 210044, Nanjing, China. Email: YujunLu.Rise@gmail.com

^b Songshan Lake Materials Laboratory, 523808, Dongguan, China. Email: kaiyang_xu@163.com

*Corresponding Author Email: 002699@nuist.edu.cn

1. User Interface of MatterMind

1.1 UI of VASP Studio




Figure S1. Home page of VASP Studio of MatterMind




Figure S2. the Part of Active VASP Run of VASP Studio




Figure S3. the Part of VASP Metrics of VASP Studio




Figure S4. the Part of VASP Metrics of VASP Studio (continued)




Figure S5. the Part of VASP Metrics of VASP Studio (continued)




Figure S6. VASP Live Logs and AI Assistant

1.2 UI of MatterGen Studio




Figure S7. Home page of MatterGen Studio of MatterMind






Figure S8. Active Run of MatterGen Studio





2 MatterGen Calculation Support Cases

2.1 Supplementary Case 1: Unconditional crystal generation with mp_20_base

To extend the checkpoint-level comparison beyond the main-text mattergen_base and space_group cases, we performed an additional unconditional generation run with mp_20_base under the same sampling setting as the baseline (batch_size = 4, num_batches = 4), yielding 16 candidate structures. The aim of this supplementary test was not to assess thermodynamic stability at the raw-generation stage, but to examine whether the MatterMind workflow remains robust and informative when applied to an alternative pretrained MatterGen checkpoint.

A qualitative overview of the mp_20_base outputs is provided in Figure S11, which displays the rendered structures of all 16 generated candidates. The corresponding structure-level metadata and screening results are listed in Table S1, including formula, atom count, symmetry, volume per atom, minimum interatomic distance, close-contact count, deduplication assignment, descriptor availability, and warning status. Together, these supplementary materials show that the batch was fully processable within the MatterMind workflow and that no trivial close-contact failure or near-duplicate repetition was observed at this sampling scale.

All 16 generated structures completed the full post-processing workflow without recorded job-level or per-structure warnings. Each frame from generated_crystals.extxyz was parsed into a pymatgen Structure, annotated with crystallographic metadata, exported as CIF, rendered as a structure image, and supplemented with Magpie and SOAP descriptors. In this sense, the mp_20_base run confirms that the workflow can consistently transform raw generative outputs into a uniformly annotated and screenable candidate set.

At the initial geometry-screening level, all 16 structures passed the short-range filter, with no close-contact event below the preset threshold of 1.5 Å. The minimum interatomic distance ranged from 1.505 Å for LaTiBe2O5 to 3.195 Å for CsNaSb. Although these observations do not support any claim of thermodynamic stability, they indicate that the generated structures were not trivially unphysical in terms of nearest-neighbor geometry and were therefore suitable for subsequent relaxation-based screening.




Figure S11. Gallery of the 16 crystal structures generated by the mp_20_base MatterGen checkpoint under unconditional sampling (batch_size = 4, num_batches = 4). Each panel shows one candidate structure after parsing from generated_crystals.extxyz and standardized rendering within the MatterMind post-processing workflow. The figure is intended to provide a qualitative overview of the structural diversity of the batch, which includes predominantly low-symmetry outputs together with a broad spread of packing density and composition. These images represent raw generated candidates prior to any further relaxation- or energy-based validation.

Table S1. Structure-level summary of the 16 candidates generated by the mp_20_base checkpoint in the supplementary unconditional sampling test. For each candidate, the table reports the reduced formula, number of atoms in the generated cell, assigned symmetry, volume per atom (V/atom), minimum interatomic distance (Min dist), number of detected close-contact events, deduplication-group assignment (Dedup), and the availability of SOAP and Magpie descriptors, together with any recorded warning message (Warn). All 16 candidates completed the workflow successfully, produced downstream artifacts, and showed zero close-contact events in the initial geometry screening.

#	Formula	Atoms	Symmetry	V/atom	Min dist	Close contacts	Dedup	SOAP	Magpie	Warn
0	CaP	8	2 P-1	25.96	2.95 (Ca-P)	0	G10 x1	159390	Yes	0
1	Sn3Ir(RhC)2	8	8 Cm	15.43	2.12 (C-Sn)	0	G1 x1	159390	Yes	0
2	Dy2HgAs	4	12 C2/m	24.59	3.16 (Dy-Hg)	0	G5 x1	159390	Yes	0
3	ZrPd4	10	1 P1	16.05	2.81 (Pd-Zr)	0	G4 x1	159390	Yes	0
4	LaTiBe2O5	9	1 P1	11.91	1.51 (Be-O)	0	G9 x1	159390	Yes	0
5	MoWSe3S	12	1 P1	33.35	2.35 (Mo-S)	0	G0 x1	159390	Yes	0
6	Sr7GeIrO2	11	1 P1	29.23	2.19 (Ge-Ir)	0	G13 x1	159390	Yes	0
7	Ti2ReNi	4	1 P1	15.1	2.68 (Ni-Ti)	0	G3 x1	159390	Yes	0
8	Tb3MgCrAgS6	12	1 P1	20.72	2.55 (Cr-S)	0	G8 x1	159390	Yes	0
9	Tl4FeClO6	12	1 P1	17.63	1.91 (Fe-O)	0	G2 x1	159390	Yes	0
10	SrMnF4	12	1 P1	14.51	2.31 (F-Mn)	0	G12 x1	159390	Yes	0
11	SmSi4	5	1 P1	20.32	2.38 (Si-Si)	0	G6 x1	159390	Yes	0
12	Rb2Tl2F5	9	1 P1	24.83	2.22 (F-Tl)	0	G14 x1	159390	Yes	0
13	CsNaSb	6	2 P-1	40.7	3.19 (Na-Sb)	0	G15 x1	159390	Yes	0
14	YbEuO2	4	2 P-1	13.8	2.37 (Eu-O)	0	G7 x1	159390	Yes	0
15	LiCoO2	12	1 P1	8.87	1.89 (Co-O)	0	G11 x1	159390	Yes	0

No obvious redundancy was observed within this batch. Using StructureMatcher with $l_{tol} = 0.2$, $s_{tol} = 0.3$, and $angle_{tol} = 5^\circ$, the 16 candidates were partitioned into 16 distinct deduplication groups. Crystallographically, low-symmetry outputs dominated the batch, with 11 structures assigned to P1, three to P-1, one to Cm, and one to C2/m. The generated set also remained compositionally diverse within this small sample, spanning binary to quinary compositions. Using raw cell volume per atom as a simple screening descriptor, the values ranged from 8.87 Å³/atom for LiCoO₂ to 40.70 Å³/atom for CsNaSb, indicating that the checkpoint sampled both compact and relatively open structural motifs.

Representative follow-up candidates can be selected from different regions of this geometry-density range. For example, CaP (P-1; 25.96 Å³/atom; 2.95 Å) and Dy₂HgAs (C2/m; 24.59 Å³/atom; 3.16 Å) provide comparatively clean starting points for relaxation, whereas LiCoO₂ should be treated more cautiously because of its unusually compact raw cell (8.87 Å³/atom) and P1 assignment. LaTiBe₂O₅ also warrants caution because its minimum distance lies only marginally above the close-contact threshold. Conversely, CsNaSb and MoWSe₃S illustrate that the checkpoint can also access relatively open candidates within the same batch.

Overall, this supplementary mp_20_base case demonstrates that MatterMind can produce a fully annotated, non-redundant, and geometry-screened candidate set when operated with an alternative MatterGen checkpoint. At the same time, the present evidence remains limited to geometry-, symmetry-, and diversity-level screening. Because no formation energies, relaxed energies, electronic properties, or other stability indicators were included in the structured metrics, no strong conclusion can be drawn regarding thermodynamic realizability at this stage.

2.2 Supplementary Case 2: Sr-V-O-conditioned crystal generation with the chemical_system checkpoint

To examine chemically conditioned generation in the MatterMind MatterGen workflow, we performed an additional experiment using the chemical_system checkpoint with $batch_size = 4$, $num_batches = 4$, $properties_to_condition_on = \{chemical_system: Sr-V-O\}$, and $diffusion_guidance_factor = 2.0$, yielding 16 candidate structures. Unlike the unconditional appendix cases, this experiment was designed not only to assess workflow robustness and geometric plausibility, but also to evaluate how effectively conditional generation steers the outputs toward the requested Sr-V-O chemical system. A qualitative overview of the generated structures is provided in Figure S12, and the corresponding structure-level screening metrics are summarized in Table S2.

All 16 candidates completed the MatterMind post-processing workflow successfully. Each generated frame was parsed from generated_crystals.extxyz, converted into a pymatgen Structure, assigned crystallographic metadata, and associated with CIF export paths, rendered structure-image paths, Magpie composition descriptors, and SOAP-based structural summaries in the structured metrics. No job-level warnings or per-structure warning lists were recorded, indicating full downstream artifact and descriptor coverage for the entire batch.

Because this run was explicitly conditioned on Sr-V-O, the most important screening quantity is condition fidelity. At the element-set level, all 16 generated structures remained within the target subset {Sr, V, O}, and no foreign elements appeared in the batch. However, exact ternary matching was not achieved for every candidate. Twelve of the 16 structures were strict Sr-V-O ternaries, whereas four omitted vanadium and collapsed into Sr-O binaries, namely SrO₂, two distinct SrO₃ structures, and SrO₈. Thus, the dominant conditioning error mode in this run was vanadium omission rather than the introduction of chemically unrelated species.

At the geometry-screening level, the batch was mixed but still largely usable. Twelve of the 16 structures showed zero close-contact events under the 1.5 Å threshold, whereas four structures accumulated a total of 16

close-contact events. Importantly, the geometric issues were concentrated in a small subset of outputs. Among the 12 exact-match Sr-V-O ternaries, 10 passed the close-contact screen cleanly. Only SrV3O14 and SrV2O9 were flagged, with minimum O-O distances of 1.197 and 1.218 Å, respectively. By contrast, the off-target Sr-O binaries accounted for most of the short-contact burden, including one SrO3 candidate with four close contacts and SrO8 with nine close contacts. As shown in Table S2, the geometrically cleaner outputs were therefore concentrated mainly within the correctly matched Sr-V-O ternary subset.




Figure S12. Rendered views of the 16 crystal structures generated by the chemical_system MatterGen checkpoint under Sr-V-O-conditioned, diffusion-guided sampling (batch_size = 4, num_batches = 4, properties_to_condition_on = {chemical_system: Sr-V-O}, diffusion_guidance_factor = 2.0). The gallery provides a qualitative overview of the generated batch, including both exact-match Sr-V-O ternaries and a small number of off-target Sr-O binaries arising from partial conditioning failure. All structures shown here are raw generated candidates after standardized parsing and rendering within the MatterMind workflow, prior to any relaxation- or energy-based validation.

Table S2. Structure-level summary of the 16 candidates generated by the chemical_system checkpoint in the Sr-V-O-conditioned supplementary experiment. For each candidate, the table reports the formula, atom count, assigned symmetry, volume per atom (V/atom), minimum interatomic distance (Min dist), number of close-contact events, deduplication-group assignment (Dedup), SOAP availability, Magpie availability, and warning status. The table shows that all 16 candidates completed the post-processing workflow, that all remained within the Sr-V-O element subset, and that the main deviation from exact conditioning was vanadium omission in four Sr-O binary outputs.

#	Formula	Atoms	Symmetry	V/atom	Min dist	Close contacts	Dedup	SOAP	Magpie	Warn
0	Sr2V3O9	14	1 P1	16.59	1.68 (O-V)	0	G4 x1	855	Yes	0
1	SrO2	3	2 P-1	15.87	1.53 (O-O)	0	G14 x1	855	Yes	0
2	Sr2VO5	8	1 P1	17.43	1.70 (O-V)	0	G10 x1	855	Yes	0
3	SrVO2	4	12 C2/m	13.58	2.27 (O-V)	0	G11 x1	855	Yes	0
4	Sr3V2O6	11	1 P1	17.82	1.88 (O-V)	0	G13 x1	855	Yes	0
5	SrV3O14	18	1 P1	19.35	1.20 (O-O)	2	G0 x1	855	Yes	1
6	SrO3	8	1 P1	16.08	1.48 (O-O)	4	G8 x1	855	Yes	1
7	Sr6VO7	14	8 Cm	17.13	2.06 (O-V)	0	G15 x1	855	Yes	0
8	Sr2V3O10	15	1 P1	19.49	1.66 (O-V)	0	G3 x1	855	Yes	0
9	SrVO3	10	4 P2_1	15.02	1.96 (O-V)	0	G6 x1	855	Yes	0
10	SrO3	4	8 Cm	19.6	1.76 (O-O)	0	G9 x1	855	Yes	0
11	SrV2O9	12	1 P1	17.51	1.22 (O-O)	1	G1 x1	855	Yes	1
12	SrVO3	10	2 P-1	15.41	1.76 (O-V)	0	G7 x1	855	Yes	0
13	SrO8	18	1 P1	16.2	1.27 (O-O)	9	G2 x1	855	Yes	1
14	SrVO2	4	8 Cm	13.59	2.23 (O-V)	0	G12 x1	855	Yes	0
15	Sr2VO7	10	12 C2/m	12.06	1.92 (O-V)	0	G5 x1	855	Yes	0

No trivial structural redundancy was observed at this sampling scale. Under the applied StructureMatcher tolerances ($\text{ltol} = 0.2$, $\text{stol} = 0.3$, and $\text{angle_tol} = 5^\circ$), all 16 generated structures were assigned to distinct deduplication groups. This is noteworthy because several reduced formulas recurred in the batch, including two SrVO_2 structures, two SrVO_3 structures, and two SrO_3 structures. The deduplication result therefore indicates that repeated stoichiometries did not collapse into repeated crystal structures; instead, the checkpoint sampled structurally distinct realizations even when nominal composition overlapped.

Crystallographically, the batch remained low-symmetry-heavy but was not exclusively so. Eight structures were assigned to $P1$, three to Cm , two to $C2/m$, two to $P-1$, and one to $P2_1$. Restricting attention to the 12 exact-match Sr-V-O ternaries gives a slightly cleaner distribution: six were assigned to $P1$, two to Cm , two to $C2/m$, one to $P-1$, and one to $P2_1$. Thus, although triclinic outputs still numerically dominated, the conditioned run also produced a meaningful subset of medium-symmetry ternary candidates, including both $C2/m$ and Cm solutions within the exact-match Sr-V-O subset. This symmetry distribution, together with the rendered structures in Figure S12, suggests that conditioned generation can yield more crystallographically regular candidates even within a small sample.

Within the geometrically clean and compositionally correct subset, several candidates emerge as practical first-pass targets for follow-up relaxation. One SrVO_2 structure crystallized in $C2/m$ with only four atoms per cell, a minimum V-O distance of 2.267 Å, and a volume per atom of 13.58 Å³/atom, making it one of the cleanest and most computationally tractable structures in the batch. A second SrVO_2 structure appeared in Cm with a very similar density (13.59 Å³/atom) and a minimum V-O distance of 2.229 Å, providing a distinct structural realization at the same stoichiometry. Sr_2VO_7 , also assigned to $C2/m$, was the densest clean exact-match ternary at 12.06 Å³/atom and showed no close-contact violations, while SrVO_3 appeared twice in distinct structural realizations, one in $P2_1$ and one in $P-1$, both passing the geometry screen. Together, these candidates illustrate that the conditioned workflow can simultaneously deliver composition fidelity, acceptable short-range geometry, and nontrivial structural diversity within the requested Sr-V-O search space.

At the same time, this batch also clarifies the practical limitation of the conditioning mode. Although the target chemistry was well respected at the element-set level, exact conditioning was imperfect, and low-symmetry outputs remained common. Moreover, several oxygen-rich candidates approached or crossed the short-contact threshold, especially within the off-target Sr-O subgroup. Accordingly, this experiment should not be interpreted as demonstrating deterministic generation of valid Sr-V-O phases. Rather, it shows that the `chemical_system` checkpoint can substantially bias generation toward the desired chemistry, while still requiring post-generation screening to separate clean target-matching ternaries from off-target or geometrically problematic outputs.

Overall, this revised Sr-V-O-conditioned case provides a clearer appendix-level demonstration of conditional materials generation within MatterMind than the earlier unconditioned example. In the present batch, the workflow confined all 16 candidates to the Sr-V-O element subset, produced 12 exact-match ternaries, and yielded 10 geometrically clean target-matching candidates suitable for downstream triage. The current evidence, however, remains limited to composition-, geometry-, diversity-, and symmetry-level screening. No formation energies, relaxed energies, oxidation-state analysis, or electronic properties were available in the present outputs, and therefore no strong claim can be made regarding thermodynamic realizability or functional behavior.

2.3 Supplementary Case 3: Property-conditioned generation with the dft_mag_density_hhi_score checkpoint

To further extend the appendix-level comparison of MatterGen checkpoints, we performed an additional guided generation experiment using the `dft_mag_density_hhi_score` checkpoint with `batch_size = 4`, `num_batches = 4`, `properties_to_condition_on = {dft_mag_density: 0.15, hhi_score: 0.8}`, and `diffusion_guidance_factor = 2.0`, yielding 16 candidate structures. Unlike the earlier unconditional and chemistry-conditioned cases, this experiment was intended to examine how the MatterMind workflow behaves when MatterGen is guided by continuous property targets rather than by a fixed chemical system. A qualitative overview of the generated structures is provided in Figure S13, and the corresponding structure-level screening metrics are summarized in Table S3.




Figure S13. Rendered views of the 16 crystal structures generated by the `dft_mag_density_hhi_score` MatterGen checkpoint under property-conditioned, diffusion-guided sampling (`batch_size = 4`, `num_batches = 4`, `properties_to_condition_on = {dft_mag_density: 0.15, hhi_score: 0.8}`, `diffusion_guidance_factor = 2.0`). The gallery provides a qualitative overview of the batch, highlighting its strong concentration in H-rich compositions and its overwhelmingly low-symmetry character. All structures shown here are raw generated candidates after standardized parsing and rendering within the MatterMind workflow, prior to any relaxation- or energy-based validation.

Table S3. Structure-level summary of the 16 candidates generated by the dft_mag_density_hhi_score checkpoint in the property-conditioned supplementary experiment. For each candidate, the table reports the formula, atom count, assigned symmetry, volume per atom (V/atom), minimum interatomic distance (Min dist), number of close-contact events, deduplication-group assignment (Dedup), SOAP availability, Magpie availability, and warning status. The table shows that all 16 candidates completed the post-processing workflow and were structurally nonredundant at the StructureMatcher level, but that only one candidate, FeH2, passed the uniform close-contact screen without warning.

#	Formula	Atoms	Symmetry	V/atom	Min dist	Close contacts	Dedup	SOAP	Magpie	Warn
0	H2O2	10	1 P1	10.2	0.95 (H-O)	6	G2 x1	3330	Yes	1
1	HO3	4	1 P1	14.32	0.94 (H-O)	2	G1 x1	3330	Yes	1
2	H2O2	6	1 P1	8.99	1.02 (H-O)	4	G3 x1	3330	Yes	1
3	H5O3	8	1 P1	10.66	0.96 (H-O)	5	G4 x1	3330	Yes	1
4	FeHO2	8	1 P1	10.29	0.94 (H-O)	2	G12 x1	3330	Yes	1
5	Fe3H8O7	18	1 P1	8.31	0.94 (H-O)	10	G10 x1	3330	Yes	1
6	H5O3	8	1 P1	7.66	1.09 (H-O)	10	G5 x1	3330	Yes	1
7	H7O4	11	1 P1	7.33	0.94 (H-O)	7	G7 x1	3330	Yes	1
8	FeH2	3	12 C2/m	7.95	1.72 (Fe-H)	0	G14 x1	3330	Yes	0
9	H5O3	16	1 P1	9.01	0.97 (H-O)	12	G6 x1	3330	Yes	1
10	LiFe(HO)2	12	1 P1	10.55	0.97 (H-O)	4	G15 x1	3330	Yes	1
11	H2O	3	1 P1	8.96	0.97 (H-O)	3	G8 x1	3330	Yes	1
12	Fe2HNO2	12	1 P1	9.4	1.00 (H-N)	2	G13 x1	3330	Yes	1
13	C2NO2	5	1 P1	11.41	1.26 (C-O)	4	G0 x1	3330	Yes	1
14	Fe2H4O5	11	1 P1	8.55	0.96 (H-O)	4	G11 x1	3330	Yes	1
15	H7N2	9	1 P1	8.78	1.04 (H-N)	7	G9 x1	3330	Yes	1

All 16 candidates completed the MatterMind post-processing workflow successfully. Each frame was parsed from generated_crystals.extxyz, converted into a pymatgen Structure, assigned crystallographic metadata, and associated with CIF export paths, rendered structure images, Magpie composition descriptors, and SOAP-based structural summaries in the structured metrics. No job-level warnings or per-structure warning lists were recorded, indicating full downstream artifact and descriptor coverage for the entire batch. Operationally, therefore, this run again confirms that the workflow can transform raw generative outputs into a structured and inspectable candidate set without manual intervention.

For this particular property-conditioned case, however, an important limitation appears immediately at the analysis stage. Although the sampling run was explicitly conditioned on `dft_mag_density` = 0.15 and `hhi_score` = 0.8, the structured metrics do not report realized values of `dft_mag_density`, `hhi_score`, or any surrogate evaluation of these target properties for the generated outputs. As a result, the present dataset does not allow direct assessment of property-conditioning fidelity. This case should therefore be interpreted primarily as a workflow- and output-distribution study rather than as evidence that the requested continuous property values were successfully matched.

Chemically, the batch was strongly concentrated in hydrogen-containing light-element systems. Fifteen of the 16 structures contained H, 14 contained O, 6 contained Fe, 3 contained N, and only one structure each contained Li or C. The compositional complexity was also narrow, with 10 binaries, 4 ternaries, and 2 quaternaries. Eight candidates were H/O-only binaries, including repeated appearances of H₂O₂ and H₅O₃-type stoichiometries. Thus, under the present property-conditioning setup, the model sampled a chemically concentrated region of composition space dominated by H-rich species rather than a broad spread of inorganic crystalline chemistries.

Despite this chemical concentration, the batch remained structurally nonredundant at the geometry-matching level. Under the applied StructureMatcher tolerances (`ltol` = 0.2, `stol` = 0.3, and `angle_tol` = 5°), all 16 structures were assigned to distinct deduplication groups. This is notable because several reduced formulas recurred in the batch, including H₅O₃ three times and H₂O₂ twice. As shown in Table S3, repeated stoichiometries therefore did not collapse into identical crystal structures, but instead appeared as distinct structural realizations.

The main weakness of this batch lies in geometric screening. Using the platform's uniform short-contact threshold of 1.5 Å, only one of the 16 structures passed cleanly with zero recorded close contacts, whereas the remaining 15 candidates accumulated 82 close-contact events in total. In coarse terms, 4 structures showed mild violations with 1 to 3 close contacts, 6 showed moderate violations with 4 to 6 contacts, and 5 showed severe violations with 7 or more contacts. The smallest minimum interatomic distance in the batch was 0.936 Å for FeHO₂, while the largest minimum-distance value was 1.719 Å for FeH₂. On a literal geometry-screening basis, this batch is therefore substantially weaker than the earlier `mattergen_base`, `mp_20_base`, and `chemical_system` cases.

At the same time, the interpretation of these warnings requires some care. The shortest-distance pair was H-O in 12 of the 16 structures, H-N in 2 structures, C-O in 1 structure, and Fe-H in 1 structure. Because this batch is overwhelmingly hydrogen-rich, many flagged short contacts arise in X-H motifs rather than in heavy-atom clashes. The uniform 1.5 Å cutoff therefore appears especially stringent for this run and may function here as a conservative warning rather than as definitive evidence that every flagged structure is chemically nonsensical. Even so, the result still carries a practical message: without relaxation or element-pair-aware thresholds, most outputs from this run cannot be treated as immediately clean crystal candidates.

Crystallographically, the batch was almost completely dominated by low-symmetry outputs. Fifteen of the 16 candidates were assigned to P1, and only one structure, FeH₂, was assigned to a higher-symmetry space group, C2/m. This is a markedly stronger low-symmetry bias than in the earlier appendix cases and suggests that, at

least in this small sample, the `dft_mag_density_hhi_score` checkpoint under the present settings did not preferentially generate crystallographically regular candidates. This low-symmetry tendency is also evident in the rendered gallery shown in Figure S13.

Within the batch, FeH₂ is the clearest first-pass follow-up target. It is the only structure with zero close contacts, the only non-P1 structure, and the one with the largest minimum interatomic distance in the batch at 1.719 Å. It is also among the most compact candidates, with a volume per atom of 7.95 Å³/atom. By contrast, most other structures are better regarded as relaxation-first cases rather than immediate screening targets. Among these, C₂N₂O₂ is notable because it is the only carbon-containing output and shows the largest minimum distance among the flagged structures at 1.261 Å, while Fe₂HNO₂ may be retained as a compositionally distinct N-containing outlier with only two recorded close contacts. Nevertheless, both remain below the geometric acceptance threshold adopted here.

Overall, this property-conditioned experiment provides a useful mixed or negative appendix-level result. The generation-and-analysis pipeline executed robustly end to end, but the resulting batch was chemically narrow, strongly low-symmetry-biased, and difficult to screen under a simple uniform geometric rule. In that sense, the value of MatterMind in this case lies not in demonstrating high-quality property-conditioned crystal generation, but in making the limitations of the run immediately visible through structured metadata, deduplication labels, minimum-distance statistics, symmetry assignments, descriptor coverage, and visualization artifacts. As seen jointly in Figure S13 and Table S3, the workflow exposes a clear contrast between operational success and scientific screening quality.

The limitation of this supplementary case is therefore more severe than in the earlier MatterGen examples. In addition to the absence of formation energies, relaxed energies, and electronic properties, even the targeted conditioning quantities themselves are not reported for the generated outputs. Consequently, no strong claim can be made about whether the requested `dft_mag_density` and `hhi_score` values were actually realized, and only a weak claim can be made about raw structural quality before relaxation. Accordingly, this case should be understood as an appendix-level demonstration showing that MatterMind can faithfully expose when a guided MatterGen run is operationally complete yet scientifically limited, thereby helping identify which checkpoints and conditioning modes warrant deeper downstream validation.

3 Ground-truth comparison

3.1 Ground-truth comparison for the VTST interpretation layer

To assess the reliability of the LLM-assisted interpretation layer for transition-state calculations, we performed a ground-truth comparison using the VTST NEB case `vtst_ground`, corresponding to the `relax_first` workflow for Sr2Ti2O5 (Table S4). In this comparison, the LLM was provided with `vtst_metrics.json`, and its final analysis output was evaluated against the raw VTST artifacts, including `neb.dat`, `spline.dat`, `exts.dat`, `nebresults.txt`, `image_energy_table.csv`, the image-level OUTCAR.gz and OSZICAR files, and the endpoint OUTCAR files. Agreement was scored on a three-level scale, with Full = 1.0, Partial = 0.5, and None = 0.0. Whenever the structured summary conflicted with the raw VTST artifacts, the raw files were treated as authoritative.

The comparison shows that the LLM performed well in workflow-level triage and in identifying the main scientific limitations of the calculation. In particular, it correctly recognized that the run, although technically completed, should not be treated as a validated saddle-point result, correctly identified image 2 as the highest-energy image, and correctly diagnosed the spline barrier as unreliable because it was strongly inflated relative to the raw NEB profile under a severely undersampled image set. It also accurately described the path as non-monotonic and structurally discontinuous, which is consistent with the raw energy sequence and the structure-matcher failures recorded between consecutive images. These results indicate that the LLM was effective at extracting the main quality-control conclusions from the structured VTST summary.

Table S4. Ground-truth comparison between the LLM-assisted interpretation and raw VTST artifacts for the relax_first NEB case of Sr₂Ti₂O₅. Agreement was scored as Full = 1.0, Partial = 0.5, and None = 0.0, with raw VTST files treated as authoritative whenever they disagreed with the generated summary JSON. The comparison shows strong LLM performance in QC triage, highest-energy image identification, and path-shape diagnosis, but weaker performance in barrier-definition precision and inferences about mechanism when the raw outputs remain ambiguous.

Module / Case	Ground-truth source	Evaluation target	Ground-truth value / conclusion	LLM interpretation	Agreement	Score	Notes
VTST, vtst_ground	job.json, vtst_metrics.json, 00-04/OUTCAR.gz	Overall run credibility	The process exited successfully (return_code = 0), but the scientific result is degraded: images 01-03 did not reach the ionic criterion before NSW = 80, so this is not a validated converged NEB path.	The run has low credibility and should not be treated as a validated saddle-point search.	Full	1.0	Good high-level triage. The model focused on scientific reliability rather than only job success.
VTST, vtst_ground	neb.dat, image_energy_table.csv, exts.dat, vtst_metrics.json	Highest-energy image / TS image	Image 2 is the highest-energy image. vtst_metrics.json also reports ts_image_index = 2.	Highest-energy image is Image 2.	Full	1.0	Exact match.
VTST, vtst_ground	neb.dat, exts.dat, image_energy_table.csv	Raw barrier definition	The raw VTST energies are 0.000000, 0.044045, 0.960280, -0.094160, -0.116292 eV relative to image 0. This means the forward barrier from the initial state is 0.960280 eV, while 1.076572 eV is the full energy span between the highest point and the lower final endpoint.	The raw barrier is about 1.08 eV.	Partial	0.5	The model repeated the vtst_metrics.json number correctly, but it used it as a generic "barrier" without separating forward barrier from full path energy span.
VTST, vtst_ground	spline.dat, neb.dat, vtst_metrics.json, INCAR_neb	Spline reliability	barrier_spline_eV = 4.189969 is far larger than the raw VTST profile, while only 3 intermediate images were used. This is a strong sign of spline overfitting / undersampling.	The spline barrier is not trustworthy and mainly reflects severe undersampling / spline distortion.	Full	1.0	Correct diagnosis and appropriate caution.

Module / Case	Ground-truth source	Evaluation target	Ground-truth value / conclusion	LLM interpretation	Agreement	Score	Notes
VTST, vtst_ground	neb.dat, structure_change_summary in vtst_metrics.json	Path shape and continuity	The path is non-monotonic, rising to image 2 and then dropping below the initial state. Consecutive structure-matcher failures appear for pairs [1, 2] and [2, 3], so the path is not structurally smooth.	The path is non-monotonic with an internal peak and significant structural discontinuities.	Full	1.0	Correct on both energy-profile shape and continuity warning.
VTST, vtst_ground	nebreults.txt, image_energy_table.csv, INCAR_neb, 00-04/OUTCAR.gz	Convergence / force diagnosis	nebreults.txt reports force values of 0.000000 for every image and even emits gnuplot warnings about missing valid force points. EDIFFG = -0.08 is loose for NEB. However, raw OUTCAR.gz shows images 00 and 04 reached required accuracy, while 01-03 did not.	The zero-force output likely reflects a parsing gap or incomplete termination; convergence is too loose, and all five images are treated as not finished.	Partial	0.5	The parser-gap explanation and loose-convergence warning are well supported. The overstatement is the claim that all five images failed; the endpoints did converge.
VTST, vtst_ground	structure_change_summary in vtst_metrics.json, endpoint_initial/OUTCAR, endpoint_final/OUTCAR	Mechanism-level interpretation	One oxygen moves by 2.611 Å; endpoint symmetry changes from P4/mmm to Pmmm; the TS image is Pm; endpoint fit is false. These facts support a large structural rearrangement, but they do not by themselves uniquely prove the exact mechanism.	The path likely reflects significant distortion / defect migration rather than a simple localized hop, and the final state may be a reconstructed configuration.	Partial	0.5	The caution is reasonable, but the mechanism wording is more specific than the raw evidence strictly requires.

The weaker points in this case were concentrated in quantitative precision and mechanism-level specificity. The LLM reproduced the larger 1.08 eV value reported in `vtst_metrics.json`, but treated it as a generic “barrier” without distinguishing the forward barrier from the initial state (0.960280 eV) from the full energy span between the highest-energy image and the lower final endpoint (1.076572 eV). It also adopted an overly strong statement that all images failed to converge, whereas the raw `OUTCAR.gz` files show that the endpoint images did reach the required accuracy and that the non-convergence problem was concentrated in the intermediate images. Similarly, although the LLM’s caution about large structural distortion was reasonable, its mechanism-level phrasing went somewhat beyond what could be established unambiguously from the raw files alone. Overall, the average agreement score for this case was 0.79, indicating that the LLM was strong at scientific triage and reliability assessment, but less precise when raw outputs remained ambiguous or when the structured summary collapsed distinct physical meanings into a single number.

An additional caveat emerged at the parser level. Although `vtst_metrics.json` reported that `exts.dat` was missing or contained no parseable points, the raw `exts.dat` file was in fact present and parseable. This discrepancy was therefore treated as a summary-layer defect rather than an LLM error. For this reason, the raw VTST files were taken as the ground-truth reference whenever they disagreed with the generated metrics summary.

3.2 Ground-truth comparison for the Fe Wannier and post-Wannier interpretation layer

To evaluate the reliability of the LLM-assisted interpretation layer in the Wannier and post-Wannier branch, we carried out a ground-truth comparison for the bcc Fe case with spin-orbit coupling (Table S5). This comparison included three connected subcases: the Wannier post-processing diagnostics, the Berry-curvature-derived anomalous Hall conductivity (AHC) analysis, and the Fermi-surface export step. The LLM inputs were the generated `wannier_metrics.json` and `postw90_metrics.json` files, whereas the ground-truth reference was taken from the raw Wannier and post-Wannier artifacts, including `wannier90.wout`, `wannier90.win`, `wannier90_centres.xyz`, `wannier90_hr.dat`, `wannier90-ahc-fermiscan.dat`, `wannier90.bxsf`, and the corresponding job configuration files. As in the VTST comparison, agreement was scored using Full = 1.0, Partial = 0.5, and None = 0.0, and raw files were treated as authoritative whenever they disagreed with the generated metrics summaries.

The Fe dataset shows the strongest overall performance among the four comparison cases, with a mean agreement score of 0.88. The LLM was particularly reliable in workflow-level interpretation, artifact triage, and in calibrated restraint about what should and should not be claimed from the available outputs. In the Wannier-post subcase, it correctly interpreted the meaning of the degraded label, correctly identified the role of the disentanglement warning, and correctly recognized that the absence of volumetric Wannier plots under SOC was expected rather than a sign of workflow failure. In the AHC subcase, it accurately described the scan configuration, the number of sampled points, the presence of three Cartesian components, and the maximum absolute AHC magnitude. In the Fermi-surface subcase, it correctly identified successful `wannier90.bxsf` export, correctly reported the Fermi energy and requested mesh settings, and appropriately refrained from overclaiming physical quantities such as pocket character, nesting, spin texture, or effective masses from the raw export artifact alone. Taken together, these results show that the LLM was particularly strong at workflow-aware interpretation and at preventing unsupported physical overstatement.

Table S5. Ground-truth comparison between the LLM-assisted interpretation and raw Wannier/post-Wannier artifacts for the bcc Fe case with spin-orbit coupling, including the Wannier-post, Berry/AHC, and Fermi-surface subcases. Agreement was scored as Full = 1.0, Partial = 0.5, and None = 0.0, with raw Wannier and post-Wannier outputs treated as authoritative whenever they disagreed with the generated metrics summaries. The results show strong LLM performance in workflow-aware interpretation, artifact-level triage, and calibrated restraint, while the main quantitative weaknesses arise from propagation of inconsistent parser-generated summary values rather than unsupported free-form invention.

Module / Case	Ground-truth source	Evaluation target	Ground-truth value / conclusion	LLM interpretation	Agreement	Score	Notes
Wannier post, Fe SOC	wannier90.wout, INCAR, wannier90.win, job.json	Overall run quality and meaning of degraded	The run completed and reached a final Wannier state. The raw warning is Maximum number of disentanglement iterations reached. LSORBIT = .TRUE. is active and both LWRITE_UNK and wannier_plot remain commented out, so missing volumetric plots are expected rather than a failure.	The Wannierisation is reliable despite a degraded status, and the degraded label comes from the disentanglement warning plus deliberately skipped volumetric plots under SOC.	Full	1.0	Strong interpretation of what the warning does and does not mean.
Wannier post, Fe SOC	wannier90.wout final-state block	Quantitative spread statistics	The final state contains 18 Wannier functions with average spread 0.9096 Å ² , max spread 1.8236 Å ² , and 6 functions above 1.5 Å ² ; none exceed 2.0 Å ² .	The model is well localized, but the answer quotes average_spread = 0.911, max_spread = 1.998, and 1215 functions above 1.5 Å ² .	Partial	0.5	The qualitative verdict is right, but several quantitative details were inherited from a faulty summary instead of the raw final-state block.
Wannier post, Fe SOC	wannier90.wout final-state block, wannier90_centres.xyz	Center localization on the Fe site	The final-state centers remain extremely close to the single Fe site. In the raw final-state block, the largest center radius is about 6.61e-4 Å.	The centers are exceptionally tight and chemically plausible, but the answer quotes a larger max_center_offset_A = 0.0108 Å.	Partial	0.5	Directionally correct, numerically inflated by the metrics layer.
Wannier post, Fe SOC	wannier90_hr.dat, file inventory in the task directory	Tight-binding model size and artifact completeness	wannier90_hr.dat reports num_wann = 18 and nrpts = 1909. The checkpoint and core interface files (.chk, .amn, .mmn, .eig, .hr.dat, .tb.dat) are present.	The answer correctly describes an 18 x 18 model with 1909 lattice points and states that the critical interface files are present.	Full	1.0	This is a clean raw-file match.
Berry / AHC, Fe SOC	wannier90.win, wannier90-ahc-fermican.dat, job.json, postw90_metrics.json	Scan configuration and peak magnitude	The run is successful. The input requests berry_task = ahc, berry_kmesh = 60 60 60, and a -0.5 to 0.5 eV scan with 0.02 eV step. The scan file contains 51 points, 3 AHC components, and a maximum absolute value of 1.979.	The answer correctly reports a successful AHC scan over 51 points and three components across a symmetric +/-0.5 eV window, with maximum absolute AHC 1.979.	Full	1.0	Quantitatively correct.

Module / Case	Ground-truth source	Evaluation target	Ground-truth value / conclusion	LLM interpretation	Agreement	Score	Notes
Berry / AHC, Fe SOC	wannier90-ahc-fermiscan.dat, wannier90.win, job.json	Interpretive restraint	The available raw outputs do not by themselves establish the physical-Fermi-level AHC, topological invariants, or k-mesh convergence. They only provide the scan itself and the chosen numerical setup.	The answer explicitly states that the current output does not establish topological character, finite-temperature transport, or numerical convergence, and that Fermi-level alignment still needs verification.	Full	1.0	Good example of calibrated restraint.
Fermi surface, Fe SOC	wannier90.win, wannier90.bxsf, job.json, postw90_metrics.json	Export success and numerical settings	The run is successful and writes wannier90.bxsf. The requested settings are fermi_energy = 5.4986400355 and fermi_surface_num_points = 120. The BXSf header stores a periodic 121 x 121 x 121 grid, consistent with a 120-interval request.	The answer correctly states that the export succeeded, reports 5.49864 eV as the Fermi energy, and describes the mesh as 120 points per dimension.	Full	1.0	The only nuance is that BXSf stores the periodic endpoint, so the file header shows 121 rather than 120.
Fermi surface, Fe SOC	wannier90.bxsf, wannier90.win, job.json	Interpretive restraint	The raw result is an isosurface export only. Pocket multiplicity, nesting, spin texture, dimensionality, and effective masses require separate visualization or analysis.	The answer explicitly says that such physical properties cannot be inferred from the current export metric alone.	Full	1.0	Again, the restraint is appropriate.

The main weakness of the LLM in this dataset was not free-form hallucination, but the propagation of incorrect quantitative values from an internally inconsistent summary layer. In the Fe Wannier-post subcase, the raw wannier90.wout final-state block reports 18 Wannier functions with an average spread of $0.9096 \text{ \AA}^2$, a maximum spread of $1.8236 \text{ \AA}^2$, and only 6 functions above $1.5 \text{ \AA}^2$, while the generated metrics summary contained inflated or inconsistent values for the maximum spread and related counts. The LLM repeated some of these faulty derived numbers, even though its qualitative conclusion that the model was localized and scientifically usable remained correct. A similar issue affected the center-localization summary, where the raw final-state centers remain extremely close to the Fe site, but the answer repeated a numerically inflated center-offset value from the metrics layer. These mismatches therefore reflect a weakness in summary-layer robustness rather than a fully independent interpretive failure by the LLM itself.

A further practical caveat concerns the post-Wannier output logging. In the Berry/AHC and Fermi-surface subcases, postw90.out was zero bytes, so execution success had to be judged from the job configuration, empty warning lists, and the presence plus contents of the generated data artifacts themselves. Even under this limitation, the LLM still produced appropriately cautious interpretations. This feature of the comparison is important because it shows that the LLM's strongest contribution in this branch is not high-precision numerical extraction under all circumstances, but robust workflow-level interpretation combined with appropriate scientific restraint.

3.3 Ground-truth comparison for the SrTiO₃ Wannier/postw90 interpretation layer

To assess the reliability of the LLM-assisted interpretation layer in the non-SOC Wannier/postw90 branch, we performed a ground-truth comparison for the cubic SrTiO₃ case, including both the Wannier-interpolated band-reconstruction task and the DOS task (Table S6). In these comparisons, the LLM inputs were the generated postw90_metrics.json files, and the corresponding final analysis outputs were evaluated against the raw post-Wannier artifacts, including wannier90-bands.dat, wannier90-path.kpt, wannier90.wpout, wannier90-dos.dat, and the job-level metadata. Agreement was again scored using Full = 1.0, Partial = 0.5, and None = 0.0, with raw files treated as authoritative whenever they disagreed with the generated metrics summaries.

Table S6. Ground-truth comparison between the LLM-assisted interpretation and raw post-Wannier artifacts for the cubic SrTiO₃ band-interpolation and DOS cases. Agreement was scored as Full = 1.0, Partial = 0.5, and None = 0.0, with raw postw90 outputs treated as authoritative whenever they disagreed with the generated metrics summaries. The comparison shows that the LLM was reliable for workflow triage, run-status checking, and non-overclaiming behavior, while its main weaknesses arose from over-trusting parser-generated summary artifacts and, in one DOS case, being more conservative than the raw output strictly required.

Module / Case	Ground-truth source	Evaluation target	Ground-truth value / conclusion	LLM interpretation	Agreement	Score	Notes
Band interpolation, SrTiO ₃	job.json, wannier90.wpout, file inventory	Execution success and artifact completeness	The task finished successfully with return_code = 0. The raw output set includes wannier90-bands.dat, wannier90-bands.gnu, wannier90-bands.py, wannier90-path.kpt, wannier90.wpout, and plots/postw90_band.png.	The task completed successfully with all standard interpolation artifacts generated and no runtime failures reported in the summary.	Full	1.0	Correct workflow-level reading.
Band interpolation, SrTiO ₃	wannier90-bands.dat, wannier90-path.kpt, wannier90.wpout, postw90_metrics.json	Band multiplicity and path sampling	Raw wannier90-bands.dat contains 14 band segments, each sampled on 756 path points. wannier90-path.kpt also reports 756 path points. The metrics JSON field n_bands = 1 and n_kpoints = 10584 is a parser artifact: 10584 = 14 x 756 is the total number of data rows, not the number of k-points.	The answer states that only 1 band was interpolated and that 10,584 k-points were sampled along the path, then treats this as an immediate configuration problem.	None	0.0	This is the main quantitative miss in this case, but it is traceable to the summary JSON rather than free-form invention.
Band interpolation, SrTiO ₃	wannier90-bands.dat, wannier90.wpout, wannier90.win	Energy window and interpretive restraint	The interpolated energy range is -1.3338227 to 11.231804 eV. The raw postw90 outputs do not include a separate DFT reference curve or an interpolation-error metric, so fidelity to the parent DFT bands cannot be established from this dataset alone.	The answer correctly reports the broad energy window and correctly withholds any claim of DFT-Wannier agreement because no explicit comparison metric is present.	Full	1.0	Appropriate restraint.
DOS, SrTiO ₃	job.json, wannier90-dos.dat, wannier90.wpout, postw90_metrics.json	Execution success and numerical consistency	The DOS task finished successfully. The setup is dos_kmesh = 20 20 20, dos_energy_min = -8, dos_energy_max = 8, dos_energy_step = 0.02. The raw DOS file contains 801 points and reaches a maximum DOS of 18.855817, matching the summary JSON.	The answer correctly identifies a successful run with 801 points over -8 to 8 eV at 0.02 eV spacing and a peak DOS near 18.86.	Full	1.0	Numerically correct.

Module / Case	Ground-truth source	Evaluation target	Ground-truth value / conclusion	LLM interpretation	Agreement	Score	Notes
DOS, SrTiO3	wannier90.wpout, wannier90-dos.dat	Availability of Fermi reference and immediate DOS-at-zero inference	wannier90.wpout explicitly reports Fermi energy (eV) = 0.000, and wannier90-dos.dat gives a finite DOS of about 4.81 at E = 0. This means a zero-energy reference is present in the raw output, and the interpolated DOS does not show a clean gap at that chosen reference.	The answer says that the Fermi position and zero-energy alignment are not reported, so metallicity or band-edge conclusions cannot be drawn.	Partial	0.5	The answer is too conservative here. The missing Fermi-reference claim is false in the raw output, but the broader caution about overinterpreting DOS without more context is still reasonable.
DOS, SrTiO3	wannier90-dos.dat, wannier90.wpout	Limits on orbital / spin interpretation	The distributed DOS data file has only 2 columns and the raw output excerpt provided here does not give orbital- or spin-resolved channels. More detailed character assignment would require additional decomposition outputs or separate analysis.	The answer correctly notes that orbital or spin decomposition is not available in the current metadata and therefore cannot support detailed character assignment.	Full	1.0	Correct limitation statement.

The comparison indicates that the LLM was generally reliable in workflow-level triage, artifact-completeness checking, and restraint against unsupported claims. In the band-interpolation subcase, it correctly recognized successful execution, the presence of the standard interpolation artifacts, and the absence of any explicit DFT-reference or interpolation-error metric from the available outputs. In the DOS subcase, it correctly identified the numerical setup, the total number of sampled energy points, and the peak DOS value, and it also correctly refrained from making orbital- or spin-resolved claims when no such decomposition was present in the raw DOS artifact. These aspects show that the LLM was effective at reading the available outputs conservatively and did not overstate what the post-Wannier files could support.

The main weakness in this dataset arose from over-trust in an incorrect parser-generated summary. In the band-interpolation subcase, the raw wannier90-bands.dat file contains 14 band segments sampled over 756 path points, while the generated summary collapsed this result into `n_bands = 1` and `n_kpoints = 10584`, where the latter is simply the total number of data rows. The LLM repeated this parser artifact and interpreted it as an immediate configuration problem. This produced the only None agreement score in the SrTiO₃ comparison and clearly illustrates a recurring pattern also seen in other cases: when the structured summary is internally flawed, the LLM may faithfully repeat the flawed summary rather than recover the correct interpretation directly from raw files that it never saw.

A related but weaker issue appears in the DOS subcase. The raw wannier90.wpout explicitly reports Fermi energy (eV) = 0.000, and the raw DOS file shows a finite DOS at that reference energy. The LLM nevertheless stated that the Fermi position and zero-energy alignment were not reported, which is too conservative relative to the raw output. Even so, its broader caution against overinterpreting metallicity or detailed band-edge character from the current DOS summary remained reasonable. This case therefore demonstrates that the LLM can fail in both directions: it may over-trust an incorrect structured summary, or it may become more conservative than the raw files strictly require.

Overall, the mean agreement score of 0.75 suggests that the LLM was useful for execution-status checking, artifact completeness, and boundary-setting on scientific claims, but less reliable when parser errors distorted the structured summary or when raw reference information was present but omitted from the generated metrics. This makes the SrTiO₃ dataset particularly informative for evaluating the interpretation layer, because it shows that the main vulnerability is not unrestricted hallucination, but dependence on the fidelity of upstream structured summaries.

3.4 Ground-truth comparison for the MatterGen interpretation layer

To evaluate the reliability of the LLM-assisted interpretation layer in the generative branch, we carried out a ground-truth comparison for a conditioned MatterGen run using the `chemical_system_energy_above_hull` checkpoint. According to the on-disk job configuration, this case was executed with `chemical_system = Au-Xe-Sb`, `energy_above_hull = 0.04`, `diffusion_guidance_factor = 2.0`, `batch_size = 2`, and `num_batches = 2`, producing a total of four generated structures (Table S7). In this comparison, the LLM was provided with `metrics.json`, and its final analysis output was evaluated against the raw generation artifacts, including `job.json`, `job.log`, `generated_crystals.extxyz`, the exported CIF files, and the associated descriptor tables. Agreement was scored using Full = 1.0, Partial = 0.5, and None = 0.0, and the raw generation files were treated as authoritative whenever they disagreed with the generated summary.

This case shows that the LLM performed well in geometry-level triage, structural diversity assessment, and epistemic restraint. It correctly recognized that all four generated structures passed the close-contact screen, correctly described the batch as nonredundant under the current deduplication settings, and correctly stated that no thermodynamic conclusion could be drawn because the available outputs contained no relaxed energies, forces, or hull values. In this sense, the model behaved appropriately when the structured metrics supported

geometric screening but did not support stronger claims regarding phase stability. The resulting mean agreement score of 0.75 indicates that the LLM was useful for summarizing what the metrics did support and for avoiding unsupported thermodynamic overinterpretation.

The main failure in this case, however, is scientifically important. Although the requested conditioning target was Au-Xe-Sb, the raw generation artifacts show that all four generated structures were in fact binary Au-Sb phases and that Xe appeared in none of the generated frames. The exact chemical-system match rate was therefore 0/4. The LLM did recognize that the batch consisted of Au-Sb structures, but it did not explicitly compare this result to the requested conditional target and therefore failed to identify the complete breakdown of chemical-system conditioning fidelity. This is the most significant error in the case, because it shows that an LLM can correctly summarize the outputs it sees while still failing to assess whether the original generation objective has actually been satisfied.

Table S7. Ground-truth comparison between the LLM-assisted interpretation and raw generation artifacts for the conditioned MatterGen case `chemical_system = Au-Xe-Sb`, `energy_above_hull = 0.04`. Agreement was scored as Full = 1.0, Partial = 0.5, and None = 0.0, with raw generation files treated as authoritative whenever they disagreed with the generated summary JSON. The comparison shows that the LLM was reliable for geometry-level screening, deduplication interpretation, and appropriate thermodynamic restraint, but failed to identify the most important result of the case: complete breakdown of the requested chemical-system conditioning fidelity.

Module / Case	Ground-truth source	Evaluation target	Ground-truth value / conclusion	LLM interpretation	Agreement	Score	Notes
MatterGen, <code>mattergen_ground</code>	<code>job.json</code> , <code>job.log</code> , <code>generated_crystals.extxyz</code>	Execution success and sample count	The generation run completed successfully with <code>return_code = 0</code> and produced 4 structures. <code>generated_crystals.extxyz</code> also contains 4 frames.	The answer consistently discusses a batch of four generated structures.	Full	1.0	Correct basic task accounting.
MatterGen, <code>mattergen_ground</code>	<code>job.json</code> , <code>generated_crystals.extxyz</code> , <code>gen_*.cif</code> , <code>magpie.csv</code>	Chemical-system conditioning fidelity	The requested chemical system is Au-Xe-Sb, but all 4/4 generated structures are binary Au-Sb phases. Xe appears in 0/4 frames, so the exact chemical-system match rate is 0/4.	The answer silently reframes the batch as "four generated Au-Sb structures" and never flags that the conditioning target Au-Xe-Sb failed completely.	None	0.0	This is the main failure in this case. The model noticed the output chemistry but did not compare it against the requested target chemistry.
MatterGen, <code>mattergen_ground</code>	<code>metrics.json</code> , <code>generated_crystals.extxyz</code> , <code>gen_*.cif</code>	Geometric validity / close-contact screening	All 4/4 structures have <code>num_close_contacts = 0</code> , empty warning lists, and minimum interatomic distances between 2.806 and 2.916 Å, well above the 1.5 Å close-contact threshold.	The answer correctly states that all four structures pass basic geometric validation and steric screening.	Full	1.0	Good geometry triage.
MatterGen, <code>mattergen_ground</code>	<code>metrics.json</code>	Deduplication / structural diversity	<code>dedup_summary.num_groups = 4</code> for 4 frames, so none of the four generated structures were merged as duplicates by the current matcher settings.	The answer correctly mentions consistent deduplication and unique groups.	Full	1.0	Correct use of the dedup summary.

Module / Case	Ground-truth source	Evaluation target	Ground-truth value / conclusion	LLM interpretation	Agreement	Score	Notes
MatterGen, mattergen_ground	job.json, metrics.json, output file inventory	Thermodynamic evidence and energy-above-hull validation	The task is conditioned on energy_above_hull, but the output bundle contains no relaxed energies, forces, or hull values for the generated candidates. Therefore the success of the energy_above_hull condition cannot be validated from the available files.	The answer correctly states that definitive thermodynamic ranking is impossible because energy / stability metrics are absent.	Full	1.0	Appropriate restraint.
MatterGen, mattergen_ground	metrics.json, generated_crystals.ext xyz, gen_*.cif	Top-K prioritization for follow-up	Structures 2 and 3 can be defended only as geometry-based screening choices: 2 has the smallest cell (4 atoms) and 3 has the lowest volume per atom. But because no candidate satisfies the requested Au-Xe-Sb chemistry and no thermodynamic data is available, these are not validated top candidates for the conditioned generation task.	The answer ranks 2 and 3 as the top two candidates with moderate-to-high confidence.	Partial	0.5	The heuristic triage is understandable, but it is not ground-truth support for the intended conditioned objective.

A secondary limitation appears in the top-K recommendation step. The LLM-selected candidates can be defended as geometry-based follow-up choices, for example in terms of smaller cell size or denser packing, but these rankings do not constitute validated recommendations for the intended conditioned-generation task because the batch already failed the target chemistry constraint and still lacks any thermodynamic evidence. In other words, the LLM was reasonably effective at internal geometric prioritization within the generated batch, but this should not be confused with success relative to the externally specified conditional objective. This distinction is important for the interpretation layer of generative workflows: geometry plausibility, conditioning fidelity, and thermodynamic support must be evaluated separately rather than collapsed into a single notion of “good” output.

An additional caveat concerns reproducibility and provenance. The raw job configuration stored on disk does not exactly match the user-provided parameter description, specifically in `batch_size`, `num_batches`, and `energy_above_hull`. For this reason, the comparison was carried out strictly against the actual on-disk artifacts rather than the user-side command description. This further reinforces the importance of grounding interpretation in raw workflow outputs rather than in intended settings alone.

4 Prompt for LLMs for Different Calculation Result Analysis

4.1 Prompt for MatterGen Calculation

SYSTEM_PROMPT = """You are the “quality reviewer for materials-generation results (MatterGen metrics.json)”. The input is the pipeline output metrics.json (ASE → pymatgen → spglib → StructureMatcher → matminer/Dscribe → (optional) CHGNet/MatGL → OVITO).

All visible output, including reasoning_content and the final answer, must be in English.

Hard rules (must be followed):

- 1) Use only the numerical values and fields in metrics.json as factual sources. It is strictly forbidden to infer properties such as “band gap / magnetism / conductivity / whether it is a semiconductor / whether it is amorphous” from experience alone. If you speculate, you must explicitly label it as “speculation / low confidence” and state which supporting calculations are missing.
- 2) A null value in tool_versions means “version not recorded / unknown”; it does not mean that the tool was not run. Whether a step produced results should be judged from the existence of corresponding output fields (such as soap_summary, render.png, relax, etc.).
- 3) You may perform derived calculations (for example, volume_per_atom = volume / natoms), but you must state the source fields and the formula used.
- 4) For each structure, output: PASS/WARN/FAIL + confidence (0–1) + evidence (cite the specific field names and values).
- 5) You must provide the “Top-K most worthwhile next-step candidates” (default K = 2), and explain the selection rationale using only evidence supported by the metrics.

Required output format:

- Part 1: One-sentence overall assessment (≤ 3 lines)
 - Part 2: Structure-by-structure review (one subsection per structure, including PASS/WARN/FAIL, key metrics, and risk points)
 - Part 3: Top-K recommendations and rationale (comparison table, including derived indicators you compute, such as volume_per_atom)
 - Part 4: Next-step action list (ranked by priority, including “fields that need additional calculations / additional recording”, such as pymatgen/dscribe/ovito version records, energy/relaxation results, etc.)
- """

CHAT_SYSTEM_PROMPT = """You are the follow-up assistant for a materials generation analysis report.

Use the existing analysis content and the conversation history as your only sources.

If the user asks for facts not present in the analysis, say what is missing and suggest rerunning the analysis.

Do not invent new material properties beyond the analysis; if you speculate, mark it as low confidence and explain why.

Answer in English only, including reasoning_content and the final answer.

Keep responses concise and actionable.

"""

4.2 Prompt for VTST NEB

VTST_SYSTEM_PROMPT = """You are a senior first-principles and transition-state analysis assistant, specialized in using VASP + VTST (NEB / CI-NEB / Dimer) to analyze diffusion paths, reaction paths, saddle points, and energy barriers.

All visible output, including reasoning_content and the final answer, must be in English.

Your task is to read the vtst_metrics.json provided by the user and produce an analysis of this VTST calculation that is oriented toward scientific and engineering decision-making, rather than simply restating the fields.

Please strictly follow the requirements below (you may adjust formatting as needed to maximize clarity for a large language model):

I. General principles

1. Base your analysis only on the input JSON. If some information is missing, explicitly say “insufficient data / cannot determine”; do not fabricate anything.
2. Prioritize judging whether the NEB/VTST result is trustworthy, and only then interpret the barrier and path.
3. The output should read like a computational diagnostic report, emphasizing:
 - whether the calculation actually ran through,
 - whether the result is reliable,
 - what the barrier is,
 - where the highest-energy image is located,
 - whether the path shape is reasonable,
 - and what should be improved next.
4. If QC flags, warnings, tails, and artifacts in the JSON contradict one another, prioritize QC fields and raw log summaries.
5. Do not overuse jargon; provide clear judgments and reasons.

II. Key analysis dimensions

Please analyze in the following order:

1. Task status and health

- whether `finished_cleanly` is true,
- whether any critical files are missing (such as `neb.dat`, `spline/exts`, or endpoint OUTCAR files required by the actual VTST mode),
- whether there are obvious warnings or errors,
- and give an overall judgment: successful / basically successful but with risks / unreliable / failed.

2. Barrier and path interpretation

- Extract the raw barrier and spline barrier.
- Identify where the highest-energy image (`ts_image_index`) is located.
- Determine whether the path is monotonically increasing, monotonically decreasing, contains an internal peak, or contains multiple peaks.
- If the highest-energy point occurs at an endpoint rather than an internal image, explicitly state: “This is more consistent with the final state being higher in energy than the initial state, rather than the true saddle point having been located.”
- If `path_monotonic` = true, explicitly explain that this usually means:
 - the final state may not be another stable state,
 - the path construction may be oversimplified,
 - or the path may not yet have crossed the true saddle point.

3. Convergence and confidence

- Examine force, energy, and convergence information for each image.
- If all forces are 0, missing, or clearly unreasonable, point out that this reduces confidence in the path.
- If the endpoints were not fully relaxed, or the number of images is too small, point that out.
- If there are electronic or ionic convergence problems, explain how they affect barrier interpretation.

4. Structural and physical interpretation

- Based on `image_table`, structure summaries, and path summaries, judge whether the process looks more like diffusion, migration, local rearrangement, configurational transformation, or merely an artificial perturbation.
- If the evidence is insufficient, say: “The current JSON does not allow a reliable determination of the specific microscopic mechanism.”

5. Next-step recommendations

Provide actionable next-step recommendations, ranked from highest to lowest priority. Recommendations should be as specific as possible, for example:

- fully relax the initial and final states separately first,
- increase the number of images,
- start with ordinary NEB and then switch to CI-NEB,
- start a Dimer calculation from the highest-energy image,
- check and tighten electronic convergence settings,
- verify that the endpoints are true local minima,
- if the path is monotonic, reconstruct a more reasonable final state or reaction coordinate.

III. Criteria to use during analysis

Use the following general criteria:

1. An “ideal NEB path” should usually have its highest point near an internal image, not at an endpoint.
2. If the path rises almost monotonically toward the final state, this usually suggests:
 - the final state is not another stable state,
 - the current path is more like pushing a stable state toward a high-energy configuration,
 - or there are too few images to cross the true saddle point.
3. If the spline barrier and raw barrier differ substantially, warn that spline fitting may amplify or distort local features.
4. If QC flags indicate missing endpoint files, abnormal forces, lack of convergence, or log anomalies, you must lower the confidence of the conclusion.
5. If the structural change is very small but the energy keeps rising, be cautious that this may only be an artificial translation/perturbation test rather than a genuine migration path.

IV. Required output format

Please strictly follow the structure below:

[Overall conclusion]

- Summarize in 2 to 4 sentences whether the calculation was successful, whether the result is trustworthy, approximately how large the barrier is, and whether the path resembles a true NEB saddle-point path.

[Key results]

- Barrier (raw / spline)
- Highest-energy image
- Path shape
- Confidence level (high / medium / low)

[Detailed diagnosis]

- Task completion status
- QC / convergence issues
- Path and barrier interpretation
- Structural / mechanism interpretation (if determinable)

[Next-step recommendations]

- Provide 3 to 6 recommendations, ranked by priority.
- Make each recommendation as specific as possible; avoid vague statements.

[One-sentence judgment]

- End with a single sentence summarizing what this result is most appropriate to be treated as, for example:
 - “A successful NEB test run, but not yet sufficient for a formal barrier report.”
 - “The endpoint and path setup are basically reasonable and already suitable as a preliminary barrier estimate.”

- “The current path looks more like an artificial perturbation than a true saddle-point search.”

V. Style requirements

- Professional, restrained, and explicit tone
- Make judgments rather than merely restating the JSON
- Do not output JSON
- Do not list raw content field by field
- Do not pretend you saw structural details that are not present in the JSON

"""

VTST_CHAT_SYSTEM_PROMPT = """You are the follow-up assistant for a VTST/NEB analysis report.

Use the existing analysis content and the conversation history as your only sources.

If the user asks for facts not present in the analysis or vtst_metrics summary, say what is missing and why it matters.

Do not invent microscopic mechanisms beyond the reported VTST evidence; if you speculate, mark it as low confidence.

Answer in English only, including reasoning_content and the final answer.

Keep responses concise, technical, and actionable.

"""

4.3 Prompt for HDF5

HDF5_VASP_SYSTEM_PROMPT = """You are a senior first-principles materials analysis assistant for VASP HDF5 pipeline outputs.

All visible output, including reasoning_content and the final answer, must be in English.

Your input is a structured HDF5 VASP metrics file produced from vaspout.h5, vasprun.xml, and related outputs.

Your job is not to restate the JSON. Your job is to evaluate run health, data reliability, and scientific usefulness for follow-up work.

Rules:

1. Use only the provided JSON as factual evidence. If a value is missing, say it is unavailable.
2. Do not invent properties such as conductivity, magnetism, topology, stability, or phonon behavior unless explicitly supported by fields in the JSON.
3. Treat qc and warnings as higher-priority evidence than inferred interpretations.
4. If tool_versions or optional sections are missing, say the data is not recorded or not available; do not assume the tool did not run.
5. Any derived quantity must cite the source fields and the calculation.

What to analyze, in order:

1. Job health and QC

- Did the run finish cleanly?
- Was electronic convergence reached?
- Were force or stress thresholds exceeded?
- Are there missing or corrupt critical files?
- Summarize reliability as High / Medium / Low.

2. Structure and crystallography

- Report formula, atom count, volume, and symmetry if available.
- Comment on whether the crystallography data look complete and internally consistent.
- If standardization or high-symmetry path data exist, explain what they enable for follow-up work.

3. Energy, force, stress, and electronics

- Summarize final total/free energy, max force, RMS force, stress tensor, Fermi level, band gap, and magnetization when present.
- If convergence is weak, explain how that limits interpretation.
- If band/DOS plugin results exist, summarize only what is explicitly present.

4. Scientific utility and next actions

- State whether this result is mainly useful as a finished result, a screening result, or a debugging result.
- Give concrete next steps such as tighter convergence, relaxation, denser k-mesh, band workflow, phonon workflow, or additional post-processing.

Required output structure:

```
## Overall conclusion
## QC and reliability
## Structure and electronic interpretation
## Recommended next actions
```

Style:

- Professional, concise, and evidence-based
 - Prefer judgment over repetition
 - Do not output JSON
 - Do not use fenced code blocks
- ```
"""
```

```
HDF5_VASP_CHAT_SYSTEM_PROMPT = """You are the follow-up assistant for a VASP HDF5 analysis report.
```

```
Use the existing analysis content and the conversation history as your only sources.
```

```
If the user asks for facts not present in the analysis or HDF5 metrics summary, say what is missing and why it matters.
```

```
Do not invent physical conclusions beyond the reported VASP evidence; if you speculate, mark it as low confidence.
```

```
Answer in English only, including reasoning_content and the final answer.
```

```
Keep responses concise, technical, and actionable.
```

```
"""
```

## 4.4 Prompt for Wannier

```
WANNIER_SYSTEM_PROMPT = """You are a senior Wannier90 analysis assistant for localized-orbital quality review and tight-binding readiness assessment.
```

```
All visible output, including reasoning_content and the final answer, must be in English.
```

Your input is a structured wannier\_metrics.json generated from a Wannier post workflow.

Your job is not to restate the JSON. Your job is to judge whether the Wannierisation is reliable, compact, and useful for downstream tight-binding or postw90 work.

Rules:

1. Use only the provided JSON as factual evidence. If a value is missing, say it is unavailable.
2. Do not invent band topology, transport, Berry curvature, magnetic behavior, or orbital character unless explicitly supported by fields in the JSON.
3. Treat quality\_assessment, warnings, llm\_summary, spread\_summary, and tight\_binding.compactness\_assessment as higher-priority evidence than speculation.
4. If the metrics are compact summaries, do not assume hidden details; refer only to the recorded preview or summary fields.
5. Treat this task as standalone. source\_step is provenance only, not evidence for additional conclusions.
6. Any derived quantity must cite the source fields and the calculation.

What to analyze, in order:

1. Run health and Wannier quality
  - Did Wannierisation converge?
  - Are key interface files and checkpoint files present?
  - What do the spread statistics imply about localization quality?
  - Summarize confidence as High / Medium / Low.

## 2. Localization and centers

- Summarize number of Wannier functions, average/max spread, and any clearly delocalized outliers.
- Use center\_summary and center-offset summaries to judge whether centers stay near chemically plausible bonding or atomic regions.
- If center offsets are large, explain why that weakens trust in the model.

## 3. Tight-binding compactness and usability

- Evaluate model dimension, dominant hopping scale, nearest/next-nearest hopping behavior, and truncation trends.
- State whether the model looks compact enough for larger-scale scanning or whether it remains diffuse.
- Use only recorded tight\_binding summaries and warnings.

## 4. Scientific utility and next actions

- State whether this result is mainly useful as:
  - a production-quality Wannier model,
  - a promising first-pass model,
  - or a debugging/intermediate result.
- Give concrete next steps such as revising projections, tightening disentanglement windows, checking frozen windows, rerunning with different num\_wann, or proceeding to postw90.

Required output structure:

```
Overall conclusion
Localization quality
Tight-binding readiness
Recommended next actions
```

Style:

- Professional, concise, and evidence-based
  - Prefer judgment over repetition
  - Do not output JSON
  - Do not use fenced code blocks
- ```
"""
```

WANNIER_CHAT_SYSTEM_PROMPT = """You are the follow-up assistant for a Wannier90 analysis report.

Use the existing analysis content and the conversation history as your only sources.

If the user asks for facts not present in the analysis or Wannier metrics summary, say what is missing and why it matters.

Do not invent physical conclusions beyond the reported Wannier evidence; if you speculate, mark it as low confidence.

Answer in English only, including reasoning_content and the final answer.

Keep responses concise, technical, and actionable.

```
"""
```

4.5 Prompt for POSTW90

POSTW90_SHARED_CHAT_SYSTEM_PROMPT = """You are the follow-up assistant for a Wannier/postw90 analysis report.

Use the existing analysis content and the conversation history as your only sources.

If the user asks for facts not present in the analysis or postw90 metrics summary, say what is missing and why it matters.

Do not invent physical conclusions beyond the reported postw90 evidence; if you speculate, mark it as low confidence.

Treat the current module as standalone; do not rely on upstream Wannier calculations unless the current analysis already discussed them.

Answer in English only, including reasoning_content and the final answer.

Keep responses concise, technical, and actionable.

```
"""
```

POSTW90_SYSTEM_PROMPTS = {
 "band_interp": """"You are a senior Wannier interpolation analysis assistant.
All visible output, including reasoning_content and the final answer, must be in English.

Input is a structured postw90_metrics.json for the Band interpolation module.
Analyze only the current module output. Do not rely on upstream Wannier results except when explicitly summarized in this JSON.

Focus on:

1. Whether the interpolation run itself succeeded and produced usable outputs.
2. What is actually available: k-point count, band count, energy window, generated plots/files.
3. Whether the result is sufficient for downstream comparison to DFT or publication-quality plotting.
4. What is missing: there is no direct DFT-vs-Wannier error metric unless explicitly present, so do not claim agreement or disagreement without evidence.

Required output structure:

```
## Overall conclusion  
## Data quality and coverage  
## What can and cannot be concluded from the interpolation  
## Recommended next actions
```

Style:

- Evidence-based, concise, technical
- Do not output JSON
- Do not use fenced code blocks

```
""",  
  "dos": """"You are a senior Wannier DOS analysis assistant.  
All visible output, including reasoning_content and the final answer, must be in English.
```

Input is a structured postw90_metrics.json for the DOS module.
Analyze only the current DOS result. Do not infer properties such as metallicity, semiconducting behavior, or orbital character unless the JSON explicitly supports them.

Focus on:

1. Whether DOS output was generated successfully.
2. Energy window, resolution, point count, and whether the DOS result looks practically usable.
3. What can be said about the DOS near the recorded energy range and what remains unknown.
4. Whether this output is sufficient for screening, plotting, or needs denser sampling / better alignment.

Required output structure:

```
## Overall conclusion  
## DOS data quality  
## Interpretable features and limitations  
## Recommended next actions
```

Style:

- Evidence-based, concise, technical
- Do not output JSON
- Do not use fenced code blocks

```
""",  
  "berry_ahc": """"You are a senior Berry-curvature and anomalous-Hall analysis assistant.  
All visible output, including reasoning_content and the final answer, must be in English.
```

Input is a structured postw90_metrics.json for the Berry / AHC module.

Analyze only the current module output. Do not claim topology, Chern character, Weyl physics, or experimentally relevant Hall behavior unless explicitly supported by the JSON.

Focus on:

1. Whether the AHC calculation succeeded and over what energy / chemical-potential window.
2. The number of components, scan range, point count, and maximum absolute AHC magnitude if present.
3. Whether the output is suitable for trend analysis or only for preliminary diagnostics.
4. What follow-up is required, such as denser k-mesh, symmetry checks, or comparison across Fermi levels.

Required output structure:

```
## Overall conclusion
## Reliability of the Berry/AHC run
## What the reported AHC scan does and does not establish
## Recommended next actions
```

Style:

- Evidence-based, concise, technical
- Do not output JSON
- Do not use fenced code blocks

```
"""
    "fermi_surface": """You are a senior Fermi-surface post-processing analysis assistant.
All visible output, including reasoning_content and the final answer, must be in English.
```

Input is a structured postw90_metrics.json for the Fermi surface module.

Analyze only the current module output. Do not infer pocket topology, carrier type, dimensionality, nesting, or spin texture unless explicitly supported by the JSON.

Focus on:

1. Whether the export succeeded and produced a usable BXSf/XSf file.
2. What file format and artifact set are available for downstream visualization.
3. Whether the result should be treated as a successful export, a partial export, or a failed visualization attempt.
4. What the user should do next to inspect or validate the Fermi surface.

Required output structure:

```
## Overall conclusion
## Export success and artifact quality
## What can and cannot be concluded from the current result
## Recommended next actions
```

Style:

- Evidence-based, concise, technical
- Do not output JSON
- Do not use fenced code blocks

```
"""
    "boltzwann": """You are a senior BoltzWann transport analysis assistant.
All visible output, including reasoning_content and the final answer, must be in English.
```

Input is a structured postw90_metrics.json for the BoltzWann transport module.

Analyze only the current transport output. Do not claim absolute experimental transport quality unless the JSON explicitly supports it.

Focus on:

1. Whether transport outputs were generated successfully.
2. Which quantities are available: conductivity, Seebeck, transport DOS, energy or temperature scan coverage, and

component count.

3. Whether the result is useful for trend analysis, relative screening, or only debugging.

4. Important caveats such as relaxation-time assumptions, missing units/context, or incomplete scan ranges.

Required output structure:

Overall conclusion

Transport data quality

What the current transport outputs support

Recommended next actions

Style:

- Evidence-based, concise, technical

- Do not output JSON

- Do not use fenced code blocks

""",

}